



Containers & Serverless

High Performance Computing (Computació d'Altes Prestacions)

Josep Lluís Berral-García – josep.ll.berral@upc.edu

Jordi Torres Viñals – jordi.torres@upc.edu



Introduction

“Encapsulate your environments and applications...
...to better replicate, migrate and deploy anywhere”



Session Objectives

- Explain the difference of containerized applications vs. VMs
- Design a container deployment schema for our AI applications
- Make a list of 5 advantages of containers vs. VMs and raw systems






Containers and Supercomputers

CONTAINERIZATION

Isolation and Environments

- Containerization of Services
 - In between “as is” and “virtualization”



Isolation and Environments

- Isolation of Services

- In between running “as is” and full system virtualization

- No HW is virtualized/emulated
 - The Host OS and HW are mostly used
 - An environment is deployed for a set of tasks → No persistence expected

- Hypervisor

- Crafted to share parts of the OS with the container instances
 - Potential security issues → No full isolation (as in VMs)
 - Reduce weight of container instances

Containers Life-Cycle

Containers ~ isolation of environments as a process, oriented to scenarios of “create-use-destroy”

- Virtual Machines
 - An instance (from an “image”) and disk volume exist per VM
 - Modifications are persistent: changes remain
- Containers
 - An “image” is to be instanced (deployed, “container”) as many times as wanted
 - When a container finishes, it is not re-used



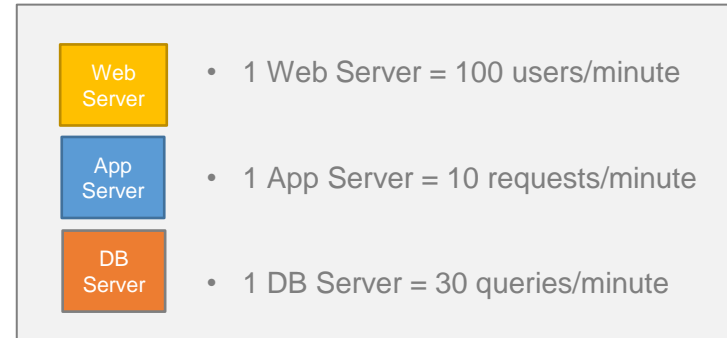

Containers Life-Cycle

Containers ~ isolation of environments as a process, oriented to scenarios of “create-use-destroy”

- Virtual Machines
 - Life-cycle:
 - A VM “instance” is created from an “image” + virtual disk volume exist
 - The “instance” (VM) is started and used
 - Changes in the VM persist
 - The “instance” can be stopped and resumed
 - The “instance” is destroyed when is not needed anymore
- Containers
 - Life-cycle:
 - An “instance” is to be instanced (deployed) as many times as wanted
 - Each “instantiation” → a “container”
 - At the end of its use, the “instance” is finished
 - When a container finishes, it is not re-used
 - Changes in the container are not stored

Containers Life-Cycle

• Example for a Web-Service



Containers Life-Cycle

- In the example:
 - We have 3 “servers”
 - Web-server (provides data in the form of “web pages”, e.g. Apache)
 - App-server (executes Application scripts, e.g. Python script)
 - DataBase-server (manages data, e.g. MySQL)
 - Users interact with the web-server, that asks for computing data to the app server, that requires data from the DB.
 - Example: Jupyter Notebook (has an interface → served by the web-server; has an application → python session running your scripts; has a DB → in this case the file-system storing notebooks and python libraries)
 - Each server can hold X clients maximum
 - We create several instances of each server, to cover capacity
 - Each server instance is identical to the others
 - All data is stored in the filesystem/database
 - When the number of clients rise → we start more instances; when decrease → we finish instances

Containers Life-Cycle

- Blackboard Exercise

- What if we have a ML pipeline?

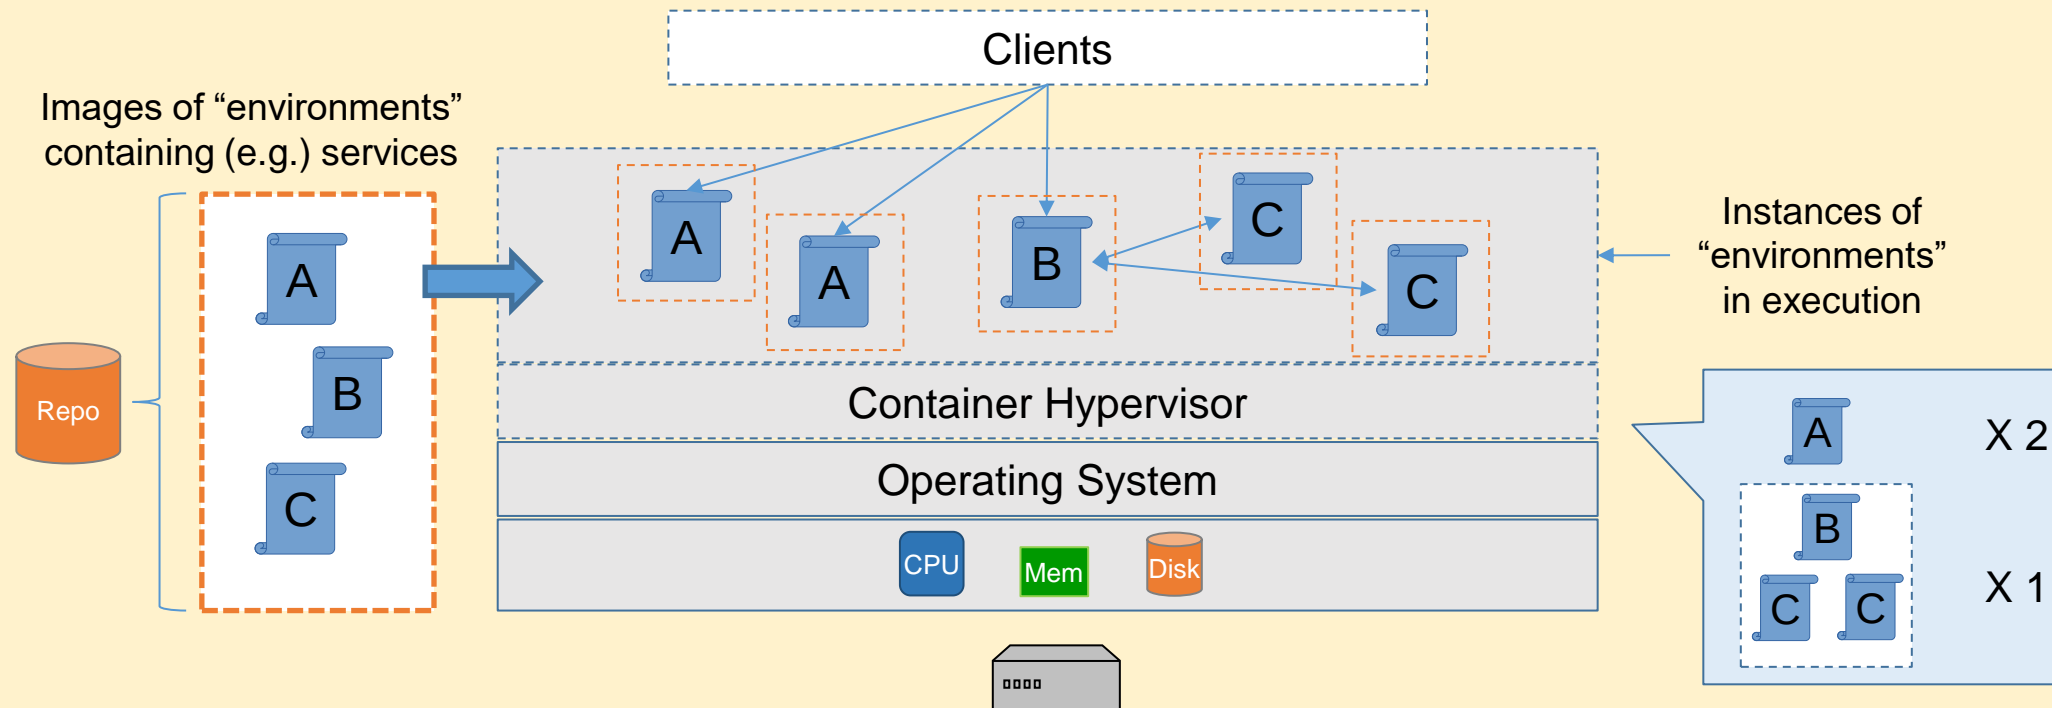
- An “app” A brings data from users
 - An “app” B trains a ML model from the last batch
 - An “app” C predicts data from A using the last ML model

- Requirements

- App A might suddenly double/triple/quadruple data input ratio N
 - App B requires X time to process a batch of size N
 - App C requires predict data A at “real-time” (maximum 1 batch in queue)

Container Architecture

- Another example of Encapsulation of Services



- Allow to deploy "instances" of services, quickly, orchestrated and in isolation

Container Architecture

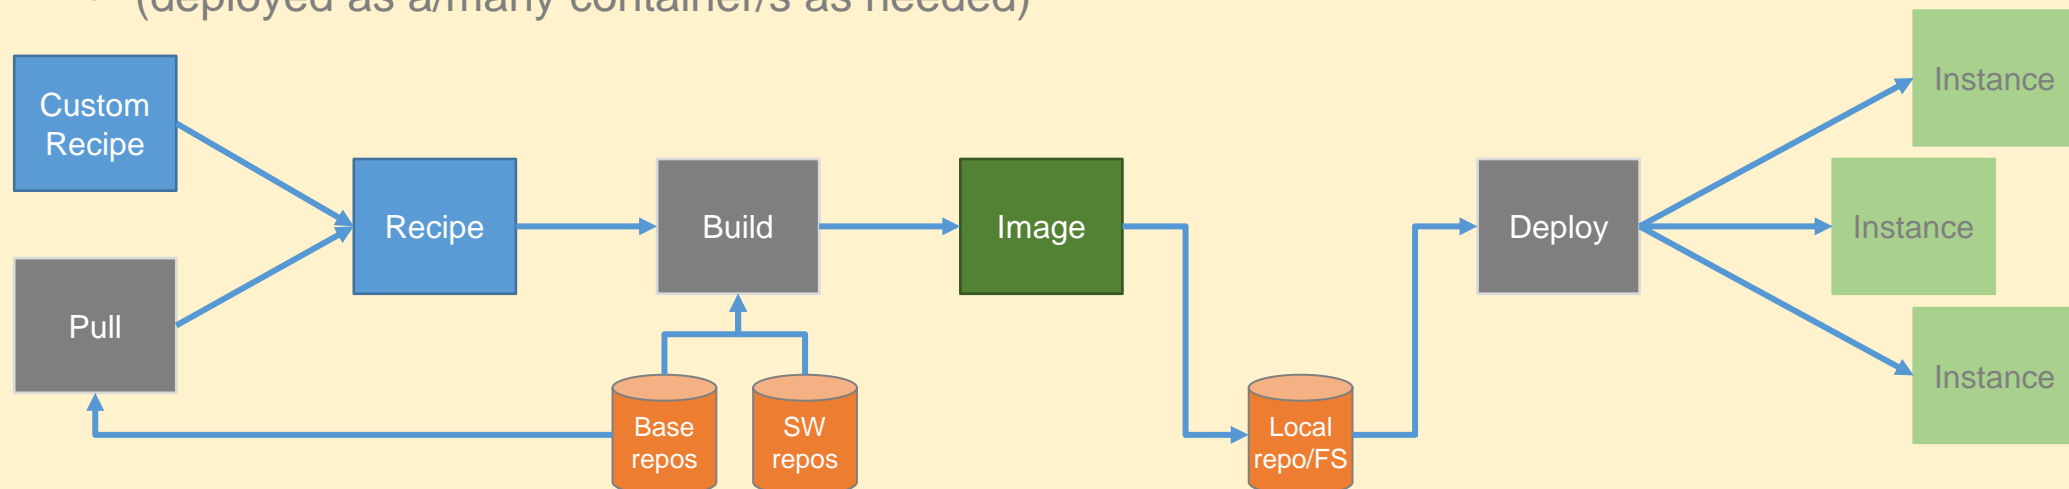
- In the example:
 - We have “images” of containers {A, B, C}
 - Each container has installed different environments and applications
 - Instances of “Container A” provide a service to the clients just alone
 - Instances of “Container B” provide a service to the clients, but require service from an instance of “Container C”
 - Request from the service Provider (person that owns A, B and C)
 - Asks que Platform/Infrastructure Provider a deployment
 - The deployment asks for
 - 2 deployments of A
 - 1 deployment containing 1 instance of B connected to 2 instances of C
 - Additionally (not in the picture):
 - The requests from clients to A are balanced between the 2 instances of A
 - The requests from B to C are balanced between the 2 instances of C
 - That way, neither A or C are overloaded → If requested, more instances of A, B and C can be instanced

Custom Containers

- The “Recipe”:
 - Instructions to “build” an image
 - “image base” + additional installs + copy files + environment variables
 - The result is an image ready to be instantiated
- Pipeline:
 - From “recipe” → build “image” of container
 - From “image” → deploy “instances” of container
 - From “image” + “receipe” → build new version of “image”

Custom Containers

- The “Recipe”:
 - Instructions to “build” an image
 - Contains the instructions to download an “image base”
 - ... plus the instructions to install additional programs
 - ... plus copy files from our system inside the container
 - ... plus do any other modification
 - E.g.: Docker → the “dockerfile”, Singularity → the “singularity file”
 - The result is an image ready to be instantiated
 - (deployed as a/many container/s as needed)





Custom Containers

- Example of Container Descriptor
 - In this case, Singularity container for SPARK and R

BootStrap: shub

From: nickjer/singularity-rstudio

%labels

Maintainer Josep Ll. Berral

Spark_Version 2.4.7

Hadoop_Version 2.7

BSC_Nord3 1.0

%help

This will run Apache Spark with an RStudio Server base, adapted to MN-IV

%runscript

exec spark-class "\${@}"

%environment

export SPARK_HOME=/nord3/spark

export PATH=\${SPARK_HOME}/bin:\${PATH}

export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-amd64

export J2REDIR=\${JAVA_HOME}/jre/

export J2SDKDIR=\${JAVA_HOME}

export JAVA_BINDIR=\${JAVA_HOME}/bin

export SDK_HOME=\${JAVA_HOME}

export JDK_HOME=\${JAVA_HOME}

export JRE_HOME=\${JAVA_HOME}/jre/

export JAVA_ROOT=\${JAVA_HOME}

%post

Software versions

export SPARK_VERSION=2.4.7

export HADOOP_VERSION=2.7

export SPARK_MIRROR=http://mirror.cc.columbia.edu/pub/software/apache/spark

Install Spark

apt-get update

apt-get install -y --no-install-recommends openjdk-8-jre

if [! -d \${SPARK_HOME}]; then

mkdir -p \${SPARK_HOME}

wget --no-verbose -O - "\${SPARK_MIRROR}/spark-\${SPARK_VERSION}/spark-\${SPARK_VERSION}-bin-hadoop\${HADOOP_VERSION}.tgz" | tar xz --strip-components=1 -C \${SPARK_HOME}

fi

MareNostrum Mount points

mkdir -p /gpfs/home

mkdir -p /gpfs/projects

mkdir -p /gpfs/archive

mkdir -p /gpfs/scratch

mkdir -p /gpfs/apps

mkdir -p /scratch

mkdir -p /.statelite/tmpfs/gpfs/projects

mkdir -p /.statelite/tmpfs/gpfs/scratch

mkdir -p /.statelite/tmpfs/gpfs/home

mkdir -p /.statelite/tmpfs/gpfs/apps/MN3

Install sparklyr: R interface for Apache Spark

Rscript -e "withCallingHandlers(install.packages(c('sparklyr'), repo= 'https://cran.rstudio.com/', clean = TRUE), warning = function(w) stop(w))"

Clean up

rm -rf /var/lib/apt/lists/*

Containers & Supercomputers

- Privileges and User-mode
 - Docker [1] → Can scale privileges (do things at root/kernel level)
 - Singularity [2] → Runs under user space and privileges
- In controlled environments...
 - ... like the MareNostrum Supercomputer...
 - ... all tasks must be contained into the user environment

[1] Docker Container: <https://www.docker.com/resources/what-container/>
[2] Singularity Container: <https://docs.sylabs.io/guides/latest/user-guide/>

Containers & Supercomputers

- Privileges and User-mode
 - Docker [1] → Can scale privileges
 - If the user instancing the container has privileges, the container can reach them
 - If the user do not have them, but the hypervisor and OS allow it, the container also can
 - This means, e.g., do things at root/kernel level
 - Singularity [2] → Runs under user space and privileges
 - The container only has the privileges of the user calling
 - And “root” is usually forbidden
- In controlled environments
 - ...like the MareNostrum Supercomputer
 - You cannot scale privileges
 - All tasks are contained into the user environment
 - We cannot use Docker in MN-IV, but we can create a Singularity container

[1] Docker Container: <https://www.docker.com/resources/what-container/>

[2] Singularity Container: <https://docs.sylabs.io/guides/latest/user-guide/>


Containers & Supercomputers

- Singularity / “User-mode” containers
 - Created (originally) for scientific environments
 - Clusters and supercomputers without privilege from users
 - Compatible with “recipes” from Docker (the most popular container hypervisor)
 - Difficult to create complex deployments automatically
 - Virtual network deployments (e.g. like in Docker-Compose or Kubernetes)
 - Access to devices that require special privileges
 - ... but good enough to deploy AI applications and science



“Pods” & Reunion of Services

- Grouping services in “Pods”
 1. Define the architecture of applications
 2. Encapsulate in “pods”
 3. Deploy in one shot

Definition of the “pod”



- Nginx → Web Server
- Flask → App Server





“Pods” & Reunion of Services

- Grouping services in “Pods”
 - Different frameworks: e.g. Kubernetes, Docker-compose, etc...
 - Infrastructure/Platform/Service described
 - A “recipe” for the whole deployment
 - Steps:
 - Define the elements (containerized services) in the “pod”
 - Indicate how they connect and how they are exposed
 - Allocate & manage available resources for the pods/containers
 - Deploy ALL in one shot

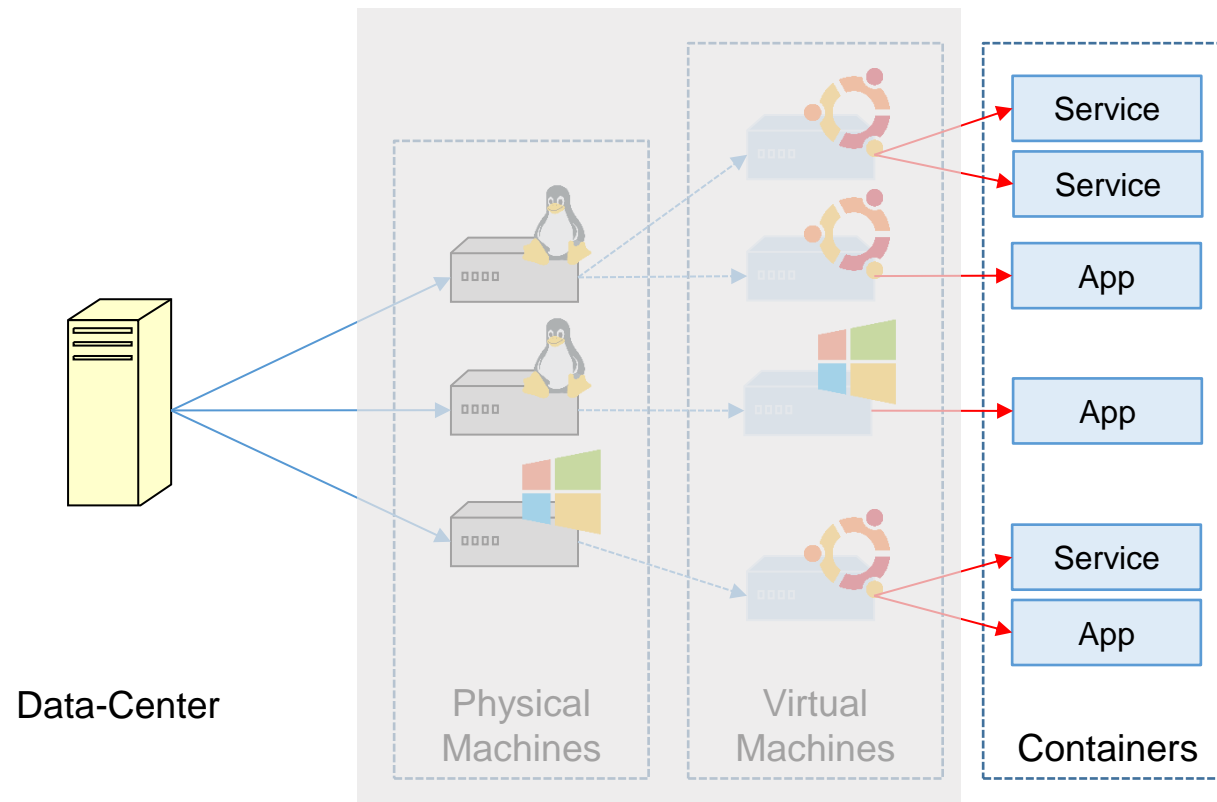
Pod description:

- 1 x Nginx → Web Server
 - Available in port 8080
- 2 x Flask → App Server
 - Application: xxxxx.py
 - Available in port 5050
- Connect
 - Nginx with Flask
 - Load balance Flask



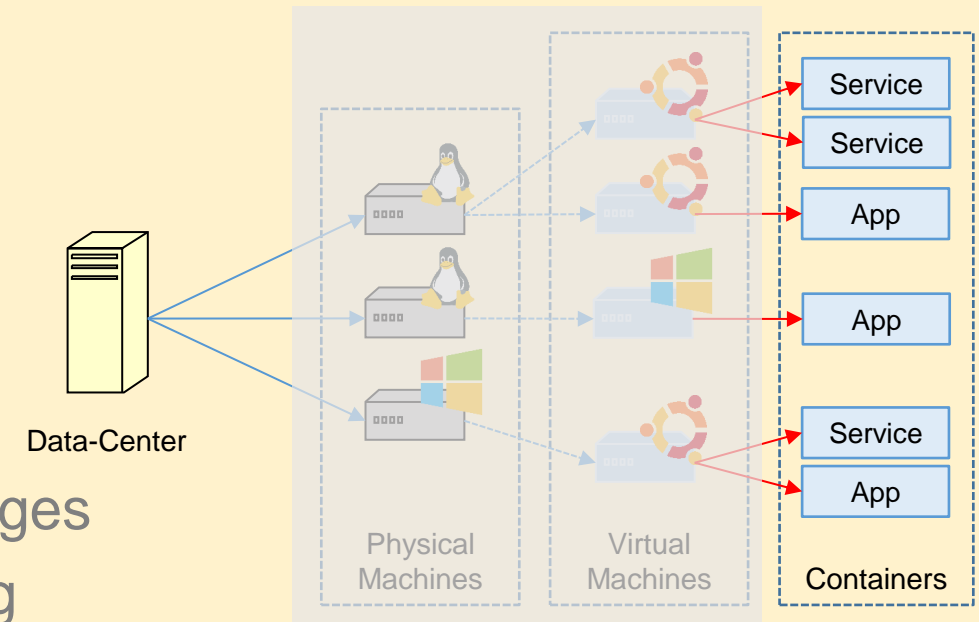
Virtualization in Resource Providers

- Resource providers can offer containers as a service
 - (inside VMs, to keep control of users and enforce security)



Virtualization in Resource Providers

- Resource providers → “Cloud owners”
 - VMs (IaaS) or Docker / Singularity / Kubernetes ... (PaaS)
 - To launch containers in scalable way
 - Runs in “transparent” platform of VMs
 - Users provide the container recipes or images
 - The manager checks the instances running
 - User does not care about what's below





Functions as a Service

SERVERLESS COMPUTING



Serverless Computing

- Serverless
 - Note: Actually... there's a “server”...
 - Micro-functions (FaaS)
 - Application logic (code) to process data
 - Common languages (Java, PHP, Python, Go, Ruby, ...)
 - Environment: prepared container with language interpreter

Serverless Computing


- Serverless
 - There's a server that receives "micro-functions"
 - Applications, mostly ETL style, that can be put to run to process massive data
 - Function as a Service
 - Micro-functions
 - User uploads the code, to process batches of data
 - The code indicates which data to retrieve and where to store the result
 - Common languages (Java, PHP, Python, Go, Ruby, ...)
 - Environment: prepared container with language interpreter
- Examples
 - Technology providers
 - AWS Lambdas
 - Google Cloud Functions, Firebase (DBs)
 - Azure Data Lake (for DBs and data-warehouses, with SQL queries)
 - Example technologies
 - OpenFaaS, Knative (Kubernetes), Lithops

Serverless Computing

- Resources consumption
 - Micro-function only consumes resources when running
 - Micro-function triggers when required only
- Scalability
 - “Infinite” scalability → Elasticity → “Deploy as many needed”
 - Not oriented for HPC, but more towards HPDA / Big-Data

Serverless Computing

- E.g., deploying μ -functions on containers from load



Serverless Computing

- Resources consumption
 - Micro-function
 - Only consumes resources when running
 - Triggers when required only, kills itself when finished
 - Has as input the data to process, and output where to store it
- Scalability
 - “Infinite” scalability → Elasticity → “Deploy as many needed”
 - Batching of data
 - Every function is just a consumer
 - Not oriented for HPC, but more towards HPDA / Big-Data
 - We can apply map-reduce techniques
- In the example
 - Depending on the load (clients or volume of data), the system triggers more or fewer micro-functions
 - For HPC / ETL applications
 - The platform deploys micro-functions depending on the batches in queue



Session Objectives

- Explain the difference of containerized applications vs. VMs
- Design a container deployment schema for our AI applications
- Make a list of 5 advantages of containers vs. VMs and raw systems



Laboratori 2 – Contenedors

PRÀCTICA DE CONTENIDORS

Laboratori

- Entorn:
 - Singularity → Hypervisor (gestiona i executa els contenidors)
 - Repositori de Singularity → Repositori públic de contenidors ja configurats
- Sistema Operatiu:
 - VM amb Ubuntu 19.04
 - Inicialment ho executarem tot dins una VM amb Singularity pre-instal·lat
 - A la pràctica tindrem 2 nivells d'imbricació (màquina → VM → contenidor → App)
 - MareNostrum-IV
 - Desplegarem el que fem a la nostra VM (contenidors) al Supercomputador
- Qüestionari:
 - Durant la pràctica cal resoldre preguntes respecte el que estem executant i observant

Laboratori

- Tenir llest l'entorn de Virtualització
 - Preparar una VM amb VBox per tenir un entorn base llest
- Instal·lar l'entorn de Contenedors
 - Instal·lar Singularity per poder executar contenidors
 - Provar a descarregar una imatge genèrica i fer-la funcionar
- Desplegar contenidors customitzats
 - Preparar una imatge “custom” que contingui pipelines de ML
 - Instanciar la imatge diverses vegades i connectar els contenidors
- Desplegar contenidors al Supercomputador
 - Copiar la imatge de contenidor creada a MareNostrum-IV
 - Preparar un llançador per a executar la pipeline de contenidors
 - Instanciar els contenidors de la pipeline com un treball de MN-IV